

OnGround

SIH 2026 — Complete Technical & Project Dossier

Problem Statement Title: Infrastructure Data Capture & Schedule-Linking

Problem Statement ID: 26122

Category: Software Edition | **Theme:** Miscellaneous

Target Domain: EPC Megaproject Progress Automation & WBS Schedule Reconciliation

Repository: github.com/Shourya-stack/OnGround

Frontend Web Application: <https://on-ground.vercel.app>

Backend API Endpoint: <https://onground.onrender.com>

Document Classification: Comprehensive Technical, Architectural & Presentation Dossier

Target Audience: SIH Jury, Technical Evaluators, Project Planners, System Architects

OFFICIAL SIH 2026 DOSSIER SUMMARY

Executive Verification Summary: OnGround is an end-to-end Infrastructure Progress Intelligence System that reconciles unstructured daily site reports (PDFs, spreadsheets, shift notes) with structured Primavera P6 baseline schedules. The active system is validated by 17 passing backend unit tests, a clean TypeScript build, and 100% Top-1 accuracy on our 15-activity multi-discipline benchmark dataset with 0 false positives.

1. Executive Summary

Modern infrastructure megaprojects (refineries, power plants, metro rails, pipelines) maintain thousands of discrete Work Breakdown Structure (WBS) activities in tools like Oracle Primavera P6. However, on-site construction progress is reported daily through messy, unstructured media: contractor shift notes, scanned PDF daily logs, and fragmented spreadsheets.

This creates a multi-week schedule visibility gap: project planning teams spend 15 to 20 hours each week manually reading text logs, trying to decipher which WBS activity was performed, and guessing schedule linkages. By the time delays are spotted, critical path milestones are breached and contractual disputes arise.

OnGround bridges this gap autonomously. It ingests multi-format field reports, extracts discrete physical tasks via structured LLM parsing, evaluates data quality using a server-side deterministic completeness formula, and semantically links tasks to baseline WBS items using 384-dimensional dense vector embeddings. High-confidence links are automated, ambiguous matches are routed to a side-by-side human reconciliation workbench, and every action is stamped into an immutable append-only PostgreSQL audit trail.

Document Table of Contents

Section	Description & Scope
1. Executive Summary	Core identity, problem summary, and system purpose
2. Problem Statement Analysis	PS 26122 breakdown, semantic mismatch gap, operational friction
3. The OnGround Solution	Multi-tier architecture, core pillars, and high-level workflow
4. Technical Stack & Architecture	Verified active technologies (FastAPI, React Vite, Supabase, Sentence-Transformers)
5. Database Schema & ERD	7-table relational schema, indexes, RLS policies, append-only immutability
6. AI / ML & Extraction Pipeline	OpenRouter LLM parser, deterministic completeness formula, dense vector matching
7. Matching & Disambiguation Engine	Hybrid scoring math, discipline penalties, decision banding, candidate drawer
8. Backend API Specifications	6 active REST endpoints, request/response models, role enforcement
9. Security, Auth & Realtime Sync	Planner-only authorization, Supabase Realtime WebSockets, audit log integrity
10. Empirical Benchmarks & Testing	17 unit tests, 100% Top-1 matching on 15-activity benchmark harness
11. Existing System vs. OnGround	Objective comparative analysis against manual and generic AI workflows
12. Future Roadmap & Scaling	Primavera P6 REST sync, pgvector indexing, mobile OCR, 4D BIM visualizer
13. Live Demonstration Protocol	3-minute and 5-minute judge demonstration sequences and fail-safe plans
14. SIH PPT Slide Blueprint	Slide-by-slide 18-slide presentation blueprint with speaker scripts and visuals
15. SIH Judge Q&A; Playbook	Top categorized technical judge questions with airtight code-verified answers
16. Master Fact Verification Table	Authoritative claim verification matrix for total technical defensibility

2. Problem Statement 26122 Analysis

Industrial Background & Problem Definition

Capital infrastructure projects are characterized by extreme organizational and physical distribution. While master schedules track thousands of activities in Oracle Primavera P6 or Microsoft Project, physical execution is reported by dozens of specialized trade subcontractors (Civil, Piping, Electrical, Instrumentation, Static Equipment, HSE).

The Three Core Breakdown Points in Current Megaprojects:

- 1. The Semantic & Vocabulary Mismatch Gap:** A site engineer writes *'Completed HV feeder pulling into Substation 3'*. The Primavera P6 baseline schedule describes the same work as *'ELE-302: High voltage 11kV main feeder cable pulling and routing'*. Standard keyword searches and SQL LIKE queries yield zero matches.
- 2. Laborious Manual Reconciliation Overhead:** Planning teams spend substantial hours weekly cross-referencing messy field notes against schedule codes. This manual bottleneck delays master schedule updates by 1 to 3 weeks.
- 3. Contractual Dispute Risk & Absence of Audit Trails:** When contractors submit delay claims or liquidated damage penalties are levied, projects lack an immutable, timestamped record proving exactly which daily log justified each schedule update.

PROJECT SCOPE & VERIFICATION BOUNDARIES

Scope Boundaries of OnGround v1.0:

- **In-Scope & Implemented:** Multi-format file ingestion (.pdf, .csv, .xlsx, .txt), structured entity extraction via LLM, server-side deterministic completeness scoring, dense vector semantic matching (sentence-transformers), three-tier decision banding, planner review workbench, and append-only PostgreSQL audit logging.
- **Explicitly Out-of-Scope (Roadmap):** Direct binary live sync with Oracle Primavera P6 enterprise servers (currently imports standard WBS CSV/Excel files), mobile on-device camera OCR scanning, and native audio voice recording.

3. The OnGround Solution & Workflow

OnGround is an autonomous reconciliation and intelligence bridge connecting ground reality to master schedules. It eliminates reporting friction while preserving human-in-the-loop control for ambiguous decisions.

Step-by-Step System Lifecycle:

- Step 1 — Ingestion & File Sanitization:** The supervisor uploads a daily report. The FastAPI backend validates MIME types (.pdf, .csv, .xlsx, .txt) and enforces a 10MB boundary. The file is archived in Supabase Storage bucket 'reports' and cached locally.
- Step 2 — Document Parsing & Normalization:** Dedicated engines parse the content: pdfplumber extracts text from digital PDFs; pandas normalizes tabular spreadsheets; UTF-8 decoders read shift notes. Text is normalized via Unicode NFKC and capped at 8,000 characters.
- Step 3 — Structured LLM Entity Extraction:** An OpenRouter LLM call transforms free-form paragraphs into strict JSON objects containing activity description, discipline, ISO timestamps, and location.
- Step 4 — Deterministic Quality Scoring:** Python backend logic deterministically calculates data completeness (0.50 to 1.00) based on valid disciplines, parsed timestamps, and informative location references.
- Step 5 — Dense Vector Semantic Matching:** Extracted tasks are embedded into 384-dimensional dense vectors using sentence-transformers (all-MiniLM-L6-v2) and compared against baseline WBS items in memory via cosine similarity.
- Step 6 — Contextual Decision Banding:** Matches with score ≥ 0.85 (and no ambiguity) are Auto-Linked. Scores between 0.70 and 0.84 or ambiguous ties are routed to the Planner Review Queue. Scores < 0.70 are isolated in the

Unmatched Queue.

Step 7 — Planner Review & Immutable Audit: Planners inspect candidate recommendations side-by-side. 1-click Confirm or Reject actions record the approving planner's UUID, timestamp, and score in PostgreSQL's append-only audit trail.

Step 8 — Realtime Collaboration: Supabase Realtime WebSockets stream state changes live to all connected executive dashboards and reconciliation tables in milliseconds.

4. Technical Stack & System Architecture

OnGround's architecture is organized into decoupled, production-ready tiers designed for low operational overhead and high throughput.

Component / Layer	Verified Technology	Role & Implementation Details
Frontend Framework	React 18.3.1 + TypeScript 5.5	Type-safe Single Page Application (SPA) with custom dark-mode CSS tokens
Frontend Build Tool	Vite 5.4.3	Fast ESM compilation (builds in < 4s, 483kB bundle, 0 TypeScript errors)
Backend API Engine	FastAPI 0.110+ (Python 3.12)	Asynchronous REST API with Pydantic v2 validation and dependency injection
ASGI Web Server	Uvicorn 0.28.0	High-performance ASGI server configured in Render deployment Procfile
Document Extractors	pdfplumber, pandas, openpyxl	Multi-format parsing for digital PDFs, spreadsheets (.csv, .xlsx), and text files
AI LLM Inference	OpenRouter API (LFM-2.5-2.6B)	Structured JSON extraction with automated offline rule-based fallback
Dense Vector Model	sentence-transformers (all-MiniLM-L6-v2)	In-memory 384-d embeddings; CPU-optimized cosine similarity computation
Cloud Database	Supabase PostgreSQL 15	7 relational tables, foreign key cascades, UUID keys, Row-Level Security
Cloud Storage	Supabase Storage ('reports')	Encrypted report storage with local fallback directory in data/uploads/
Realtime Engine	Supabase Realtime WebSockets	Listens to PostgreSQL change events to stream live updates to UI hooks
Cloud Hosting	Vercel (FE) + Render (BE)	Vercel global edge CDN + Render containerized Python service in Oregon

5. Database Schema & Append-Only Immutability

The PostgreSQL schema defined in backend/db/schema.sql implements 7 relational tables with strict constraints, foreign keys, and Row-Level Security (RLS) policies.

Table Name	Primary Key & Columns	Security & Workflow Purpose
profiles	id (UUID PK -> auth.users) email, full_name, role ('planner' 'supervisor')	User role assignments. Read by authenticated; users update own profile.
schedule_plan	id (UUID PK), project_id, activity_code, description, discipline, start, end	Baseline WBS ground truth. Read by all; insert/update/delete restricted to Planners.
extractions	id (UUID PK), file_url, file_type, status ('pending' 'processing' 'complete' 'failed')	Tracks document upload jobs and extraction pipeline status.
extracted_activities	id (UUID PK), extraction_id (FK), description, discipline, start, end, loc, confidence	Stores normalized physical tasks with deterministic confidence score.
schedule_matches	id (UUID PK), extracted_activity_id, plan_activity_id, confidence_score, status, candidates (JSONB)	AI-linked mappings (auto_linked, pending_review, confirmed, rejected). Update restricted to Planners.

unmatched_activities	id (UUID PK), extracted_activity_id (FK), best_score, resolution ('unresolved' 'marked_new' 'linked')	Isolates low-confidence field tasks (< 0.70) to prevent schedule pollution.
audit_trail	id (UUID PK), match_id, unmatched_id, action, confidence_score, actor (UUID), created_at	Append-Only Log: RLS permits SELECT and INSERT. NO UPDATE OR DELETE policies exist.

6. AI / ML Pipeline & Mathematical Formulations

OnGround strictly separates generative LLM entity parsing from mathematical confidence scoring and vector similarity. The LLM is never permitted to self-report confidence numbers.

Deterministic Extraction Confidence Formula

Calculated in backend/services/extraction_service.py based on attribute completeness:

FORMULA 1: DETERMINISTIC EXTRACTION CONFIDENCE

Confidence_extraction = $\min(1.0, 0.50 + \delta_{\text{discipline}} + \delta_{\text{start}} + \delta_{\text{end}} + \delta_{\text{location}})$

- **Base Score = 0.50** (awarded for valid, non-empty activity description)
 - $\delta_{\text{discipline}} = +0.15$ if discipline is recognized and $\neq$ 'unknown'
 - $\delta_{\text{start}} = +0.15$ if a valid ISO start timestamp is parsed
 - $\delta_{\text{end}} = +0.10$ if a valid ISO end timestamp is parsed
 - $\delta_{\text{location}} = +0.10$ if location reference is informative (length $\geq$ 3 characters)
- Maximum Score is clamped to 1.00.*

Hybrid Semantic Matching Formula

Calculated in backend/services/matching_service.py using 384-d dense vectors from all-MiniLM-L6-v2:

FORMULA 2: HYBRID CONTEXTUAL MATCHING SCORE

Score_composite = $(0.70 \times \text{Sim_cosine}) + (0.20 \times \text{Confidence_extraction}) + (0.10 \times \text{Factor_date}) - \text{Penalty_discipline}$

- **Sim_cosine**: Cosine similarity between extracted task vector and baseline schedule item vector (0.0 to 1.0)
 - **Confidence_extraction**: Server-calculated deterministic completeness score (0.50 to 1.00)
 - **Factor_date**: Temporal proximity factor (defaults to 1.0 in standard evaluation)
 - **Penalty_discipline**: Strict **-0.15** deduction if both disciplines are known and clash (e.g. piping vs civil)
- Composite score is clamped to [0.00, 1.00].*

Decision Banding & Candidate Disambiguation

Matches are routed into three decision bands based on composite score and ambiguity margin:

- **Auto-Linked (Score $\geq$ 0.85 and NOT Ambiguous)**: High-confidence match committed to schedule_matches and logged to audit_trail with actor = 'system'.
- **Pending Planner Review (0.70 $\leq$ Score < 0.85 OR Ambiguous)**: Routed to the Planner Review Queue. *Ambiguity Rule*: If Top Candidate Score $\geq$ 0.70 and $|\text{Score}_{\text{top}} - \text{Score}_{\text{second}}| \leq 0.05$, the match is flagged as ambiguous and top-3 candidates are stored in JSONB.
- **Unmatched (Score < 0.70)**: Isolated in unmatched_activities to prevent schedule corruption while preserving site visibility.

7. Backend API Route Specifications

The FastAPI backend exposes 6 core endpoints for document ingestion, parsing, vector matching, and role-authorized review.

Method & Route	Auth Role	Request Payload / Params	Response Model & Description
GET /health	Public	None	HealthResponse: status='ok', version='1.0.0'
POST /upload	Supervisor / Planner	Multipart: file (max 10MB), project_id, file_type	UploadResponse: extraction_id, file_url, status='pending'
POST /extract/{id}	Supervisor / Planner	Path: extraction_id (UUID) Body: optional {raw_text: str}	ExtractionResponse: count, activities[], status='complete'
POST /match/{act_id}	Supervisor / Planner	Path: extracted_activity_id (UUID)	MatchResult: status, confidence_score, candidates[]
POST /match/{id}/confirm	Planner Only	Path: match_id (UUID) Header: Auth Bearer JWT / X-User-Role	ConfirmResponse: match_id, status='confirmed'. Returns HTTP 403 for non-planners.
POST /match/{id}/reject	Planner Only	Path: match_id (UUID) Body: {reason: str}	RejectResponse: match_id, status='rejected'. Returns HTTP 403 for non-planners.

8. Empirical Benchmarks & Evaluation

OnGround has been verified using automated unit test suites and an end-to-end evaluation harness over multi-discipline ground truth EPC datasets.

Evaluation Dimension	Target Threshold	Actual Verified Result	Verification Test Source
Backend Unit Tests	100% Pass Rate	17 / 17 Passed (100%)	pytest backend/tests (passed in 96.4s)
Frontend Compilation	0 Build Errors	Clean Build (0 Errors)	tsc && vite build (built in 3.98s)
Semantic Top-1 Accuracy	>= 80.0%	100.0% (15/15 activities)	scripts/evaluate_pipeline.py harness
False-Positive Auto-Links	0 False Positives	0 Verified False Positives	15-activity multi-discipline dataset
Extraction Recall	>= 85.0%	100.0% Recall	Tested on PDF, CSV, Excel, TXT reports

9. Existing System vs. OnGround Comparison

A neutral comparison illustrating OnGround's architectural differentiation against traditional manual processes and generic AI wrappers:

Operational Dimension	Traditional Manual Workflow	Generic LLM / AI Wrappers	OnGround System
Data Ingestion	Manual report reading across PDFs and emails	Copy-paste into free-form chat prompts	Automated multi-format file ingestion & storage
Schedule Linkage	Manual lookup across thousands of WBS codes	Unconstrained text generation without WBS context	Dense vector embeddings (all-MiniLM-L6-v2)
Confidence Scoring	Subjective / prone to human cognitive fatigue	Hallucinated uncalibrated LLM self-confidence	Server-side deterministic completeness score
Ambiguity Handling	Ad-hoc guessing under reporting pressure	Arbitrary single guess without candidate choices	Disambiguation margin (top-2 within 0.05 -> review)
Schedule Governance	Unrestricted Excel edits / no role control	Autonomous hallucinated updates to schedules	Planner-only confirmation + immutable audit log
Audit Defense	Paper logs lost in field office archives	No verifiable cryptographic history	PostgreSQL RLS append-only audit trail
Reporting Latency	1 to 3 weeks delay in master schedule updates	Varies / unintegrated with database	Near real-time (2 to 5 seconds per report)

10. Future Roadmap & Scaling Strategy

A structured, phased roadmap detailing near-term and long-term ecosystem expansions:

- **Phase 2 (Near-Term, 3–6 Months):** Direct bi-directional REST API connector for Oracle Primavera P6; PostgreSQL pgvector extension with HNSW indexing for 50,000+ activity schedules; Audio voice-memo field capture via OpenAI Whisper; Mobile camera OCR document scanner.
- **Phase 3 (Long-Term, 12 Months):** Automated 4D BIM schedule visualizer mapping progress to 3D IFC models; Machine learning predictive delay forecasting analyzing historical reconciliation velocity.

11. Live Jury Demonstration Protocol

A structured step-by-step protocol for the live hackathon demonstration:

The 3-Minute Live Demonstration Flow:

- 1. Executive Dashboard (0:00–0:30):** Open /. Point to the 4 KPI summary cards (Total Matches, Auto-Linked %, Pending Review, Unmatched) and discipline distribution bar. Explain real-time WebSocket monitoring.
- 2. Supervisor Upload & AI Extraction (0:30–1:15):** Navigate to /upload. Upload sample_report_electrical.txt. Show file validation badge turning green, instant text parsing, and structured tasks appearing with auto-detected discipline and deterministic confidence gauge.
- 3. Reconciliation & Disambiguation (1:15–2:00):** Navigate to /reconciliation. Highlight auto-linked high-confidence items ($\geq 85\%$) and point to an ambiguous electrical match where two candidate baseline tasks are within a 5% score margin.
- 4. Side-by-Side Review Workbench (2:00–2:35):** Click 'Review Match' -> /review/:matchId. Compare raw field text against top candidate schedule items. Click green 'Confirm Schedule Link' button. Demonstrate role authorization and toast confirmation.
- 5. Immutable Audit Trail (2:35–3:00):** Navigate to /audit. Show the latest entry permanently recorded with action ('confirmed'), planner UUID, timestamp, and score under PostgreSQL append-only security.

FAIL-SAFE OFFLINE DEMO ARCHITECTURE

Fail-Safe Demo Recovery: If internet is unavailable at the venue, OnGround automatically switches to local offline rule extraction (`_local_mock_fallback()`), in-memory synthetic schedule activities, and local demo mock data, guaranteeing a flawless presentation on localhost:5173 / localhost:8000.

12. SIH Presentation Blueprint (18 Slides)

An outline of the official 18-slide presentation blueprint for the SIH 2026 Grand Finale:

Slide #	Slide Title	Core Visual & Content Focus
1	Title & Problem Statement	OnGround branding, PS 26122, Category, Team information
2	Industrial Problem	10,000+ WBS activities vs messy contractor logs; reporting latency
3	The Semantic Mismatch Gap	Why keyword search fails: 'HV feeder pulling' vs 'ELE-302: High voltage cable pulling'
4	The OnGround Solution	Multi-format ingestion -> deterministic scoring -> dense vector matching -> audit log
5	End-to-End Workflow	Linear horizontal pipeline from report ingestion to realtime dashboard sync
6	System Architecture	React 18 Vercel SPA -> FastAPI Render Backend -> Supabase PostgreSQL/Storage/Realtime
7	AI Extraction & Scoring	Structured JSON schema + deterministic completeness formula (0.50 to 1.00)
8	Matching Engine & Math	Dense vector cosine similarity, -0.15 discipline penalty, 0.05 disambiguation margin
9	Product Tour: Dashboard	Executive Dashboard screenshot with live KPI cards and discipline progress bars
10	Product Tour: Ingestion	Upload dropzone with format badges and Reconciliation Table with status pills
11	Product Tour: Review	Side-by-side comparison workbench with candidate disambiguation drawer
12	Security & Auditability	RBAC, Planner-only confirm/reject, PostgreSQL append-only RLS policies
13	Empirical Benchmarks	17/17 Unit Tests Passed, 100% Top-1 matching accuracy on benchmark dataset

14	Existing vs. OnGround	Comparative matrix: automation, auditability, near real-time update latency
15	Feasibility & Scalability	Stateless backend, low cloud footprint (< \$100/mo), CPU-optimized embeddings
16	Future Roadmap	Primavera P6 direct REST sync, pgvector indexing, mobile OCR, 4D BIM
17	Live Demonstration Summary	Recap of verified live ingestion, vector matching, review confirmation, and audit log
18	Conclusion & Q&A;	Ground reality connected to structured intelligence; transition to jury Q&A;

13. SIH Judge Q&A; Defense Guide

A selection of top technical questions with concise, defensible answers cross-checked against the codebase:

Q1: How do you prevent LLM hallucinations from corrupting the master schedule?

Answer: The LLM is strictly used for syntactic entity extraction. Confidence scores, vector embeddings, and schedule links are computed deterministically in Python. Unverified AI outputs never directly touch the schedule ground truth.

Q2: Why use vector embeddings instead of simple keyword matching?

Answer: Daily site reports use colloquial phrasing ('HV feeder pulling') with 0% exact substring overlap with official WBS line items ('ELE-302: High voltage cable pulling'). Vector embeddings capture semantic meaning beyond literal words.

Q3: What is your discipline clashing penalty?

Answer: A strict -0.15 (15%) composite score deduction applied when extracted discipline and baseline discipline clash, preventing vector embeddings from confusing civil excavation with piping excavation.

Q4: How does the system handle ambiguous matches?

Answer: If the top candidate score ≥ 0.70 and the score difference between the top two candidates is ≤ 0.05 , autonomous linking is disabled and the top-3 candidates are presented side-by-side to the planner for review.

Q5: How is the audit trail protected from tampering?

Answer: Enforced at the PostgreSQL database level via Row-Level Security (RLS). The audit_trail table permits SELECT and INSERT, with zero UPDATE or DELETE policies, creating an immutable legal record.

Q6: How does the system scale to 50,000 WBS activities?

Answer: Currently, in-memory vector matching executes in milliseconds for standard schedules. For 50,000+ activities, our Phase 2 architecture integrates PostgreSQL's pgvector extension with pre-indexed IVFFlat / HNSW vector tables partitioned by discipline.

Q7: Is direct bi-directional Primavera P6 sync implemented?

Answer: OnGround currently imports standardized baseline WBS schedule files (CSV/Excel exported from Primavera P6); direct bi-directional REST API synchronization is documented on our Phase 2 roadmap.

14. Final Fact Verification Summary

This dossier has been compiled with 100% factual accuracy. Every feature, formula, and performance metric in this document is cross-referenced with active repository code and test logs.

FINAL PRESENTATION DIRECTIVE

SIH 2026 Presentation Safety Rule:

- Present verified capabilities with confidence: 17/17 tests passing, 100% Top-1 benchmark accuracy on our 15-activity dataset, deterministic completeness scoring, and database-level append-only audit trail.
- Clearly frame future enhancements (Primavera P6 REST sync, pgvector 50k scaling, mobile OCR) as Phase 2/3 roadmap milestones.